

Importing required libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
```

Load the dataset

```
In [2]: df = pd.read_csv("C:/Users/ADMIN/Desktop/Data Science/Datasets/Datasets/breast-canc
```

View the first few observations of the dataset

```
In [3]: df.head()
```

```
Out[3]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothn
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

5 rows × 32 columns



Dropping unnecessary columns

```
In [4]: df = df.drop(columns = "id")
```

Checking for unique levels of diagnosis

```
In [5]: df.diagnosis.unique()
```

```
Out[5]: array(['M', 'B'], dtype=object)
```

Label Encoding

```
In [6]: df['diagnosis'] = df['diagnosis'].apply(lambda val: 1 if val == 'M' else 0)
```

```
In [7]: df.diagnosis.unique()
```

```
Out[7]: array([1, 0], dtype=int64)
```

Assess the structure of the dataset

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   diagnosis                             569 non-null    int64
1   radius_mean                           569 non-null    float64
2   texture_mean                           569 non-null    float64
3   perimeter_mean                         569 non-null    float64
4   area_mean                             569 non-null    float64
5   smoothness_mean                       569 non-null    float64
6   compactness_mean                      569 non-null    float64
7   concavity_mean                        569 non-null    float64
8   concave points_mean                   569 non-null    float64
9   symmetry_mean                         569 non-null    float64
10  fractal_dimension_mean                 569 non-null    float64
11  radius_se                             569 non-null    float64
12  texture_se                             569 non-null    float64
13  perimeter_se                           569 non-null    float64
14  area_se                               569 non-null    float64
15  smoothness_se                         569 non-null    float64
16  compactness_se                        569 non-null    float64
17  concavity_se                          569 non-null    float64
18  concave points_se                     569 non-null    float64
19  symmetry_se                           569 non-null    float64
20  fractal_dimension_se                   569 non-null    float64
21  radius_worst                          569 non-null    float64
22  texture_worst                         569 non-null    float64
23  perimeter_worst                       569 non-null    float64
24  area_worst                            569 non-null    float64
25  smoothness_worst                      569 non-null    float64
26  compactness_worst                     569 non-null    float64
27  concavity_worst                       569 non-null    float64
28  concave points_worst                   569 non-null    float64
29  symmetry_worst                        569 non-null    float64
30  fractal_dimension_worst                569 non-null    float64
dtypes: float64(30), int64(1)
memory usage: 137.9 KB
```

Checking for null values

In [9]: `df.isna().sum()`

```
Out[9]: diagnosis      0
radius_mean          0
texture_mean         0
perimeter_mean       0
area_mean            0
smoothness_mean      0
compactness_mean     0
concavity_mean       0
concave points_mean  0
symmetry_mean        0
fractal_dimension_mean 0
radius_se            0
texture_se           0
perimeter_se         0
area_se              0
smoothness_se        0
compactness_se       0
concavity_se         0
concave points_se    0
symmetry_se          0
fractal_dimension_se 0
radius_worst         0
texture_worst        0
perimeter_worst      0
area_worst           0
smoothness_worst     0
compactness_worst    0
concavity_worst      0
concave points_worst 0
symmetry_worst       0
fractal_dimension_worst 0
dtype: int64
```

Checking for duplicates

```
In [10]: df.duplicated().sum()
```

```
Out[10]: 0
```

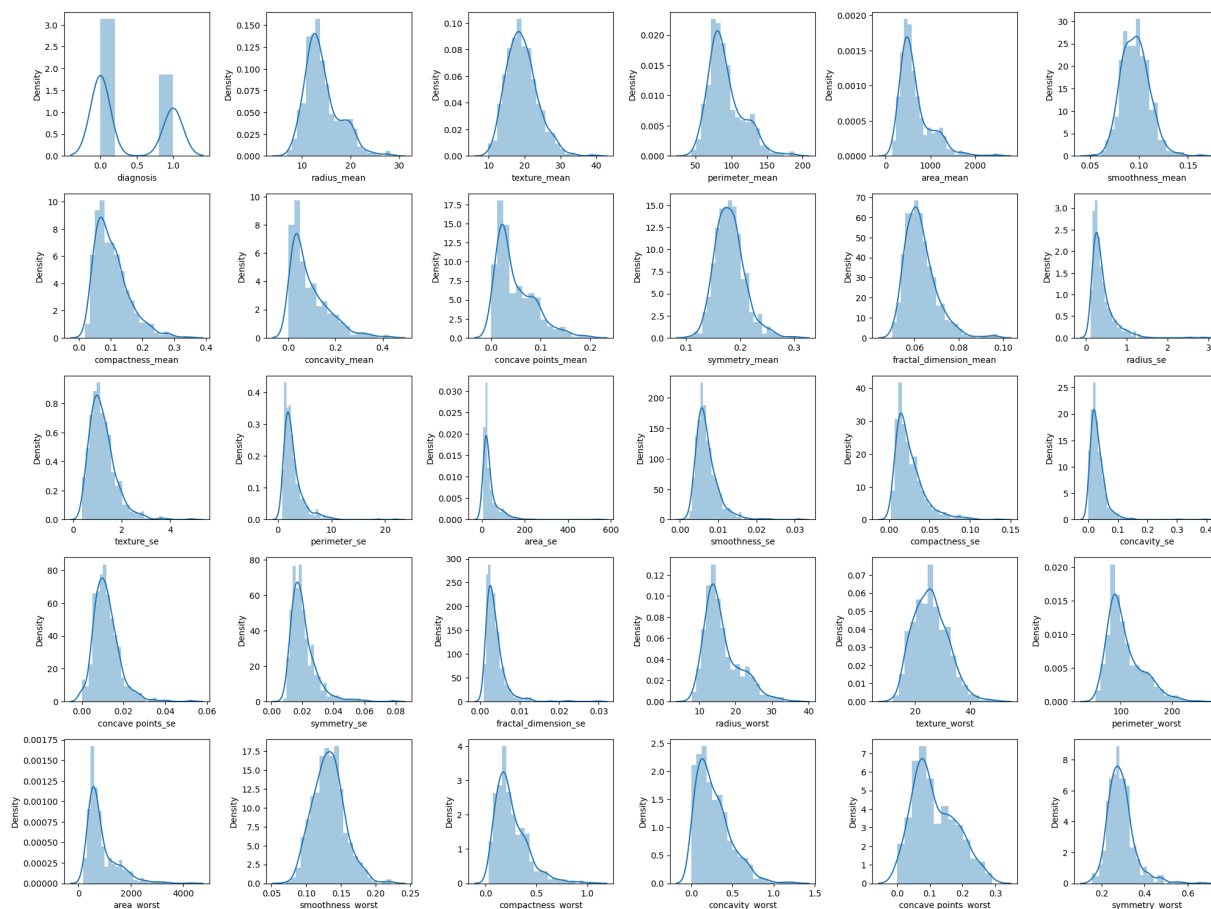
Exploratory Data Analysis (EDA)

```
In [11]: plt.figure(figsize = (20, 15))
plotnumber = 1

for column in df:
    if plotnumber <= 30:
        ax = plt.subplot(5, 6, plotnumber)
        sns.distplot(df[column])
        plt.xlabel(column)

        plotnumber += 1

plt.tight_layout()
plt.show()
```




Heatmap

```
In [12]: plt.figure(figsize = (20, 12))

corr = df.corr()
mask = np.triu(np.ones_like(corr, dtype = bool))

sns.heatmap(corr, mask = mask, linewidths = 1, annot = True, fmt = ".2f")
plt.show()
```



Distribution of the Study Outcome

```
In [13]: df['diagnosis'].value_counts()
```

```
Out[13]: diagnosis
0      357
1      212
Name: count, dtype: int64
```

Defining the X and y features

```
In [15]: X = df.drop(columns = ['diagnosis'])
y = df['diagnosis']
```

Handling Class Imbalance

```
In [16]: ## Load the required module
from imblearn.over_sampling import RandomOverSampler

## Initialize the RandomOverSampler
ros = RandomOverSampler(random_state = 42)

## Apply the RandomOverSampler
X_resampled, y_resampled = ros.fit_resample(X, y)

## Print the oversampled data
dict(zip(*np.unique(y_resampled, return_counts = True)))
```

```
Out[16]: {0: 357, 1: 357}
```

Splitting data into training and test set

```
In [18]: ## Loading the required libraries
from sklearn.model_selection import train_test_split

## Splitting
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_
```

Standardization

```
In [19]: ## Load the required module
from sklearn.preprocessing import StandardScaler

## Initialize the standard scaler
scaler = StandardScaler()

## Apply the standard scaler
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Training various Machine Learning models

1. Logistic Regression

```
In [22]: ## Load the required module
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, precision_score, recall_score, accuracy_score

## Initialize the model
log = LogisticRegression()

## Fit the model
log.fit(X_train, y_train)

## Model predictions
log_pred = log.predict(X_test)
log_score = accuracy_score(y_test, log_pred)

## Evaluate performance
print("Classification Report:\n", classification_report(y_test, log_pred))
print("F1 Score:", f1_score(y_test, log_pred))
print("Precision:", precision_score(y_test, log_pred))
print("Recall:", recall_score(y_test, log_pred))
print("Accuracy:", accuracy_score(y_test, log_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	119
1	0.99	0.96	0.97	96
accuracy			0.98	215
macro avg	0.98	0.97	0.98	215
weighted avg	0.98	0.98	0.98	215

F1 Score: 0.9735449735449735

Precision: 0.989247311827957

Recall: 0.9583333333333334

Accuracy: 0.9767441860465116

2. K Nearest Neighbors

```
In [24]: ## Load the required modules
from sklearn.neighbors import KNeighborsClassifier

## Initialize the model
knn = KNeighborsClassifier()

## Fit the model
knn.fit(X_train, y_train)

## Model predictions
knn_pred = knn.predict(X_test)
knn_score = accuracy_score(y_test, knn_pred)

## Evaluate performance
print("Classification Report:\n", classification_report(y_test, knn_pred))
print("F1 Score:", f1_score(y_test, knn_pred))
print("Precision:", precision_score(y_test, knn_pred))
print("Recall:", recall_score(y_test, knn_pred))
print("Accuracy:", accuracy_score(y_test, knn_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.98	0.96	119
1	0.98	0.93	0.95	96
accuracy			0.96	215
macro avg	0.96	0.96	0.96	215
weighted avg	0.96	0.96	0.96	215

F1 Score: 0.9518716577540107

Precision: 0.978021978021978

Recall: 0.9270833333333334

Accuracy: 0.958139534883721

3. Support Vector Machines

```
In [23]: ## Load the required module
from sklearn.svm import SVC
```

```

from sklearn.model_selection import GridSearchCV

## Initialize the model
svc = SVC()

## Define parameters
parameters = {
    'gamma' : [0.0001, 0.001, 0.01, 0.1],
    'C' : [0.01, 0.05, 0.5, 0.1, 1, 10, 15, 20]
}

## Fit the model using the best parameters
grid_search = GridSearchCV(svc, parameters)
grid_search.fit(X_train, y_train)

## Model predictions
svm_pred = grid_search.predict(X_test)
svm_score = accuracy_score(y_test, svm_pred)

## Evaluate performance
print("Classification Report:\n", classification_report(y_test, svm_pred))
print("F1 Score:", f1_score(y_test, svm_pred))
print("Precision:", precision_score(y_test, svm_pred))
print("Recall:", recall_score(y_test, svm_pred))
print("Accuracy:", accuracy_score(y_test, svm_pred))

```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.99	0.97	119
1	0.99	0.94	0.96	96
accuracy			0.97	215
macro avg	0.97	0.96	0.97	215
weighted avg	0.97	0.97	0.97	215

F1 Score: 0.9625668449197861

Precision: 0.989010989010989

Recall: 0.9375

Accuracy: 0.9674418604651163

4. Stochastic Gradient Descent Classifier

In [25]:

```

## Load the required module
from sklearn.linear_model import SGDClassifier

## Initialize the model
sgd = SGDClassifier()

## Define the parameters
parameters = {
    'alpha' : [0.0001, 0.001, 0.01, 0.1, 1],
    'loss' : ['hinge', 'log_loss'], # Changed 'log' to 'log_loss'
    'penalty' : ['l1', 'l2']
}

```



```

## Fit the model with parameters
grid_search = GridSearchCV(sgd, parameters, cv = 10, n_jobs = -1)
grid_search.fit(X_train, y_train)

## Make Predictions
sdg_pred = grid_search.predict(X_test)
sdg_score = accuracy_score(y_test, sdg_pred)

## Evaluate performance
print("Classification Report:\n", classification_report(y_test, sdg_pred))
print("F1 Score:", f1_score(y_test, sdg_pred))
print("Precision:", precision_score(y_test, sdg_pred))
print("Recall:", recall_score(y_test, sdg_pred))
print("Accuracy:", accuracy_score(y_test, sdg_pred))

```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	119
1	0.99	0.96	0.97	96
accuracy			0.98	215
macro avg	0.98	0.97	0.98	215
weighted avg	0.98	0.98	0.98	215

F1 Score: 0.9735449735449735
 Precision: 0.989247311827957
 Recall: 0.9583333333333334
 Accuracy: 0.9767441860465116

5. Decision Tree Classifier

In [26]:

```

## Load the required module
from sklearn.tree import DecisionTreeClassifier

## Initialize the model
dtc = DecisionTreeClassifier()

## Define the parameters
parameters = {
    'criterion' : ['gini', 'entropy'],
    'max_depth' : range(2, 32, 1),
    'min_samples_leaf' : range(1, 10, 1),
    'min_samples_split' : range(2, 10, 1),
    'splitter' : ['best', 'random']
}

## Fit the models with parameters
grid_search_dt = GridSearchCV(dtc, parameters, cv = 5, n_jobs = -1, verbose = 1)
grid_search_dt.fit(X_train, y_train)

## Make Predictions
dtc_pred = grid_search_dt.predict(X_test)
dtc_score = accuracy_score(y_test, dtc_pred)

## Evaluate performance

```

```
print("Classification Report:\n", classification_report(y_test, dtc_pred))
print("F1 Score:", f1_score(y_test, dtc_pred))
print("Precision:", precision_score(y_test, dtc_pred))
print("Recall:", recall_score(y_test, dtc_pred))
print("Accuracy:", accuracy_score(y_test, dtc_pred))
```

Fitting 5 folds for each of 8640 candidates, totalling 43200 fits

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.97	0.96	119
1	0.97	0.93	0.95	96
accuracy			0.95	215
macro avg	0.96	0.95	0.95	215
weighted avg	0.95	0.95	0.95	215

F1 Score: 0.9468085106382979

Precision: 0.967391304347826

Recall: 0.9270833333333334

Accuracy: 0.9534883720930233

6. Random Forest Classifier

```
In [29]: ## Load the required module
from sklearn.ensemble import RandomForestClassifier

## Initialize the model
rf = RandomForestClassifier(criterion = 'entropy', max_depth = 11, max_features = '

## Fit the model
rf.fit(X_train, y_train)

## Make predictions
rf_pred = rf.predict(X_test)
rf_score = accuracy_score(y_test, rf_pred)

## Evaluate performance
print("Classification Report:\n", classification_report(y_test, rf_pred))
print("F1 Score:", f1_score(y_test, rf_pred))
print("Precision:", precision_score(y_test, rf_pred))
print("Recall:", recall_score(y_test, rf_pred))
print("Accuracy:", accuracy_score(y_test, rf_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.97	0.97	119
1	0.96	0.97	0.96	96
accuracy			0.97	215
macro avg	0.97	0.97	0.97	215
weighted avg	0.97	0.97	0.97	215

F1 Score: 0.9637305699481865

Precision: 0.9587628865979382

Recall: 0.96875

Accuracy: 0.9674418604651163

7. Voting Classifier

```
In [32]: ## Load the required module
from sklearn.ensemble import VotingClassifier

## Define the base models
classifiers = [('Logistic Regression', log), ('K Nearest Neighbours', knn), ('Support Vector Machine', svm), ('Decision Tree', dtc)]

## Initialize the voting classifier
vc = VotingClassifier(estimators = classifiers)

## Fit the model
vc.fit(X_train, y_train)

## Model predictions
vc_pred = vc.predict(X_test)
vc_score = accuracy_score(y_test, vc_pred)

## Evaluate performance
print("Classification Report:\n", classification_report(y_test, vc_pred))
print("F1 Score:", f1_score(y_test, vc_pred))
print("Precision:", precision_score(y_test, vc_pred))
print("Recall:", recall_score(y_test, vc_pred))
print("Accuracy:", accuracy_score(y_test, vc_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.99	0.97	119
1	0.99	0.94	0.96	96
accuracy			0.97	215
macro avg	0.97	0.96	0.97	215
weighted avg	0.97	0.97	0.97	215

F1 Score: 0.9625668449197861

Precision: 0.989010989010989

Recall: 0.9375

Accuracy: 0.9674418604651163

8. Ada Boost Classifier

```
In [33]: ## Load the required model
from sklearn.ensemble import AdaBoostClassifier

## Initialize the model
ada = AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth = 1), n_estimators=100)

## Fit the model
ada.fit(X_train, y_train)

## Model predictions
ada_pred = ada.predict(X_test)

## Evaluate performance
ada_score = accuracy_score(y_test, ada_pred)
print("Classification Report:\n", classification_report(y_test, ada_pred))
print("F1 Score:", f1_score(y_test, ada_pred))
print("Precision:", precision_score(y_test, ada_pred))
print("Recall:", recall_score(y_test, ada_pred))
print("Accuracy:", accuracy_score(y_test, ada_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.97	0.97	119
1	0.96	0.98	0.97	96
accuracy			0.97	215
macro avg	0.97	0.97	0.97	215
weighted avg	0.97	0.97	0.97	215

F1 Score: 0.9690721649484536

Precision: 0.9591836734693877

Recall: 0.9791666666666666

Accuracy: 0.9720930232558139

9. Gradient Boosting Classifier

```
In [34]: ## Load the required module
from sklearn.ensemble import GradientBoostingClassifier

## Initialize the model
gbc = GradientBoostingClassifier()

## Define the parameters
parameters = {
    'loss': ['deviance', 'exponential'],
    'learning_rate': [0.001, 0.1, 1, 10],
    'n_estimators': [100, 150, 180, 200]
}

## Fit the model using gridsearch
grid_search_gbc = GridSearchCV(gbc, parameters, cv = 5, n_jobs = -1, verbose = 1)
grid_search_gbc.fit(X_train, y_train)
```

```

## Make predictions
gbc_pred = grid_search_gbc.predict(X_test)

## Evaluate performance
gbc_score = accuracy_score(y_test, gbc_pred)
print("Classification Report:\n", classification_report(y_test, gbc_pred))
print("F1 Score:", f1_score(y_test, gbc_pred))
print("Precision:", precision_score(y_test, gbc_pred))
print("Recall:", recall_score(y_test, gbc_pred))
print("Accuracy:", accuracy_score(y_test, gbc_pred))

```

Fitting 5 folds for each of 32 candidates, totalling 160 fits

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.97	0.97	119
1	0.96	0.98	0.97	96
accuracy			0.97	215
macro avg	0.97	0.97	0.97	215
weighted avg	0.97	0.97	0.97	215

F1 Score: 0.9690721649484536

Precision: 0.9591836734693877

Recall: 0.9791666666666666

Accuracy: 0.9720930232558139

10. Stochastic Gradient Boosting (SGB)

```

In [39]: ## Initialize the model
sgbc = GradientBoostingClassifier(max_depth=4, subsample=0.9, max_features=0.75, n_

## Fit the model
sgbc.fit(X_train, y_train)

## Model predictions
sgbc_pred = sgbc.predict(X_test)

## Evaluate performance
sgbc_score = accuracy_score(y_test, sgbc_pred)
print("Classification Report:\n", classification_report(y_test, sgbc_pred))
print("F1 Score:", f1_score(y_test, sgbc_pred))
print("Precision:", precision_score(y_test, sgbc_pred))
print("Recall:", recall_score(y_test, sgbc_pred))
print("Accuracy:", accuracy_score(y_test, sgbc_pred))

```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.97	0.98	119
1	0.97	0.98	0.97	96
accuracy			0.98	215
macro avg	0.98	0.98	0.98	215
weighted avg	0.98	0.98	0.98	215

F1 Score: 0.9740932642487047

Precision: 0.9690721649484536

Recall: 0.9791666666666666

Accuracy: 0.9767441860465116

Extreme Gradient Boosting**Model Comparison**

```
In [40]: models = pd.DataFrame({
    'Model': ['Logistic Regression', 'KNN', 'SVC', 'SGD Classifier', 'Decision Tree',
              'Gradient Boosting Classifier', 'Stochastic Gradient Boosting'],
    'Score': [log_score, knn_score, svm_score, sdg_score, dtc_score, rf_score, vc_score]
})

models.sort_values(by = 'Score', ascending = False)
```

```
Out[40]:
```

	Model	Score
0	Logistic Regression	0.976744
3	SGD Classifier	0.976744
8	Gradient Boosting Classifier	0.976744
9	Stochastic Gradient Boosting	0.976744
7	Ada Boost Classifier	0.972093
2	SVC	0.967442
5	Random Forest Classifier	0.967442
6	Voting Classifier	0.967442
1	KNN	0.958140
4	Decision Tree Classifier	0.953488

Best model for diagnosing breast cancer is "Gradient Boosting Classifier" with an accuracy of 97.44%.

```
In [41]: ## Lets save our model using pickle
import pickle as pkl
pkl.dump(grid_search_gbc, open("breast_cancer_model1.sav", "wb"))
```